

Machine Learning Project - Prediction of barbell lifts

Bruno BERREHUEL*

Abstract

Using devices such as Jawbone Up, Nike FuelBand, and Fitbit it is now possible to collect a large amount of data about personal activity relatively inexpensively. One thing that people regularly do is quantify how much of a particular activity they do, but they rarely quantify how well they do it. This report predicts the manner in which people did the exercises, based on data from accelerometers on the belt, forearm, arm, and dumbell. Despite that all variables with NAs have been removed, the models accuracies are very good, as shown in table 4. In this case, variables with NA aren't necessary to fit a good prediction model. Further fine tuning could be done on the model, but this will not be seen in this report.

A good model should have good accuracy for few computation time. So I recommand using the Bagged Tree model as shown in figure 4. It is possible to predict and survey, not only if people do exercices, or how many exercices they do, but most of it how well they do them.

Keywords

Machine Learning — Predictive Modeling — Monitored fit exercices

*Corresponding author: <https://fr.linkedin.com/in/brunoberrehuel>

Contents

Introduction	1
1 Dealing with datas	1
1.1 Exploratory Analysis	1
1.2 Dealing with NA's	1
2 Fitting a predictive model	2
2.1 Global quick review	2
2.2 Best 5 models study	3
3 And the winner is...	3
Acknowledgments	4
A Code	4
A.1 Load Datas	4
A.2 Simplify dataset and deal with NAs	4
NA distribution • NA mean	
A.3 Fitting a model	4
Functions to automate process • Model quick review • Best Five models	
A.4 And the Winner is...	5

Introduction

Using devices such as Jawbone Up, Nike FuelBand, and Fitbit it is now possible to collect a large amount of data about personal activity relatively inexpensively. One thing that people regularly do is quantify how much of a particular activity they do, but they rarely quantify how well they do it.

For this purpose, six candidates were asked to perform barbell lifts correctly and incorrectly in 5 different ways. Results in

the datasets are classified in 5 classes, according to how the exercises are done¹ :

- Class A : exactly according to the specification
- Class B : throwing the elbows to the front
- Class C : lifting the dumbbell only halfway
- Class D : lowering the dumbbell only halfway
- Class E : throwing the hips to the front.

1. Dealing with datas

1.1 Exploratory Analysis

There are 160 potentials predictors, for 19622 observations. After a look on the structures of the data, the following columns can be removed because they have no added value for predictive model purpose : V1, user_name, timestamps and windows columns.

There are still 153 potentials predictors, and it seems to have several NAs in the dataset, and they have to be managed.

1.2 Dealing with NA's

In this dataset, there is no classified case without NA's. So it is necessary to view the distribution of the NAs in the predictors, here the columns of the dataset, and find a solution to deal with them.

min	mean	max	sd
0.9793	0.9811	1.0000	0.0050

Table 1. Summary for mean of NA in predictors with NA

¹Source : <http://groupware.les.inf.puc-rio.br/har>

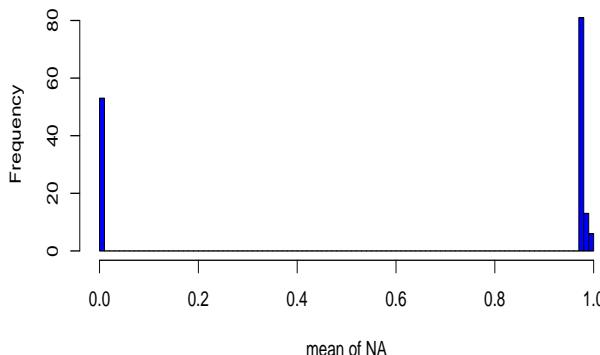


Figure 1. Distribution of NA in predictors

As seen on figure 1 and the table 1, there are only two cases for predictors with NA :

- there is no NA in the predictor
- there are a lot of NA in the predictor, more than 97.9%.

Predictors with NAs can be safely removed in order to reduce noise without available informations.

There are now 53 predictors, which will be a little easier to manage with fitting model.

2. Fitting a predictive model

The model analysis will be² :

1. test variety of models on a small sample of the dataset, with the default parameters. The train will be checked by cross-validation of 10 folds, repeated 3 times, and the distribution of the accuracy, with statistique test, will be reviewed.
2. select the bests 5 of them for accuracy, then train and test them with cross-validation on the complete dataset, under the same testing conditions.
3. finally choose the best one.

Further analysis could be done to tune the selected model in order to increase performance in accuracy and computation time, but it's not the goal of this report.

2.1 Global quick review

The Machine Learning algorithms are choosen to represent a good mix of algorithms. For this classification case, the 21 models chosen are listed in table 2.

Bagged Flexible Discriminant Analysis
C5.0
Stacked AutoEncoder Deep Neural Network
Multivariate Adaptive Regression Spline
Stochastic Gradient Boosting
Generalized Partial Least Squares
High Dimensional Discriminant Analysis
Linear Discriminant Analysis
k-Nearest Neighbors
Mixture Discriminant Analysis
Naive Bayes
Neural Network
Parallel Random Forest
Nearest Shrunken Centroids
Penalized Discriminant Analysis
Random Forest
CART
Stabilized Nearest Neighbor Classifier
Support Vector Machines with Linear Kernel
Support Vector Machines with Radial Basis Function Kernel
Bagged CART

Table 2. Models reviewed

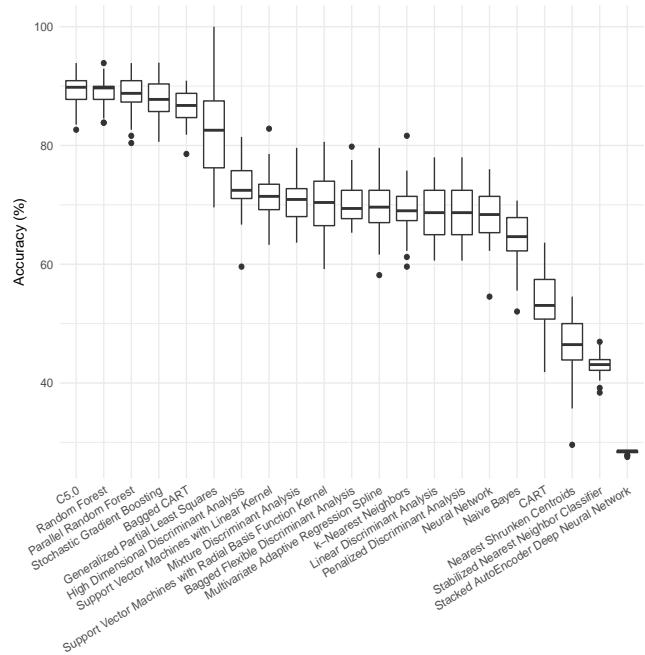


Figure 2. Accuracy of each model on quickview dataset

As seen on figure 2, the five bests models for accuracy are :

- C5.0
- Random Forest
- Parallel Random Forest
- Stochastic Gradient Boosting
- Bagged CART

²Source : <http://machinelearningmastery.com/evaluate-machine-learning-algorithms-with-r/>

The models have pretty good accuracy, from 89.1% to 86.4%. The *Bagging CART model* is especially interesting because of its small computation time, 9 seconds, in comparaison of the 82 seconds for the C5.0 model³. The statistiques tests for the means are significant, as we fail to reject that the trues means are not equals to the samples means, so the samples are representatives of the populations. The 95% confidences intervals are presented in table 3.

C5.0	88.1094	90.1804
Random Forest	87.9019	89.7806
Parallel Random Forest	87.2972	89.7675
Stochastic Gradient Boosting	86.5983	88.9134
Bagged CART	85.3466	87.5238

Table 3. Confidences intervals for models accuracy

2.2 Best 5 models study

The 5 bests models are trained on the full dataset, with cross-validation with 10 folders, repeated 3 times. Then the distribution of the accuracies are ordered and boxplotted.

	Accuracy	Time
C5.0	99.66	1455.00
Random Forest	99.53	2099.00
Parallel Random Forest	99.50	1043.00
Bagged CART	98.86	165.00
Stochastic Gradient Boosting	96.37	656.00

Table 4. Accuracy and computation time for the 5 bests models

The most important result is that the accuracies are very, very good, around 99%, and with small variations around the mean, as seen with the boxplots on the figure 3.

Note the small decrease in performance of the Stochastic Gradient Boosting model, which is now worse than the Bagging CART model.

The computations times have a very broad set, from 1 times 13 depending on the model considered.

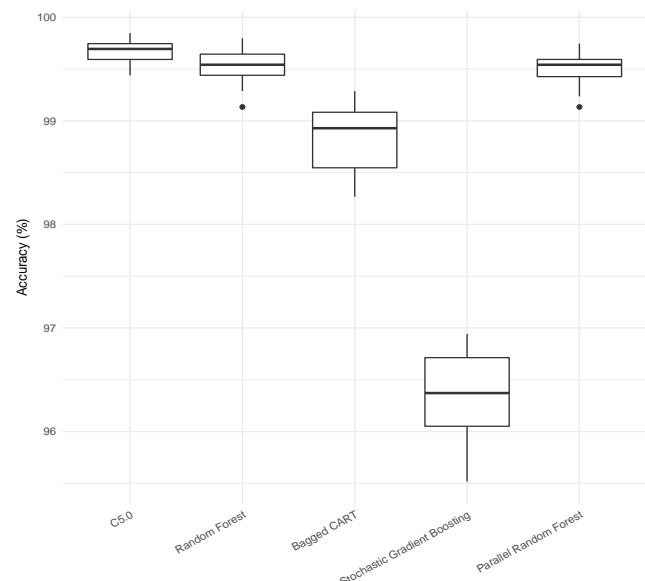


Figure 3. Accuracy of the five best models on total dataset

3. And the winner is...

A good predictive model should be very accurate and very fast.

All the four finalists models are very accurate, so the difference will be in computation time, and the clear winner is here the *Bagged CART Model*, with only 165 seconds of computation, in comparaison of 2099 seconds of the Random Forest, which is the slowest model. The difference appears more clearly on the figure 4.

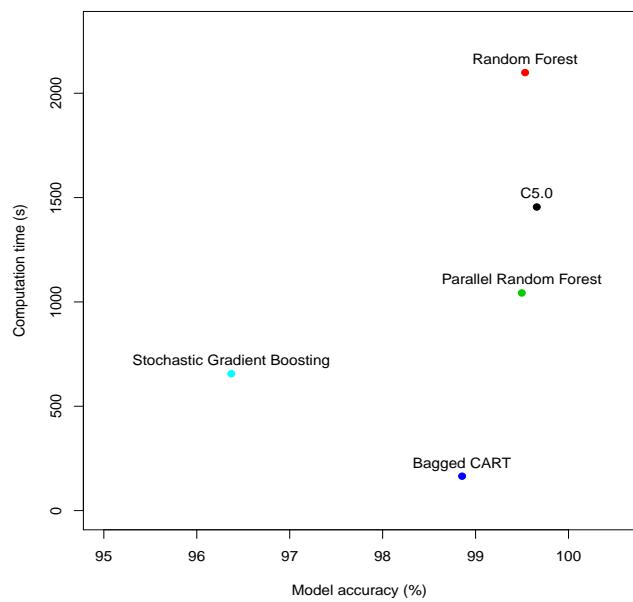


Figure 4. Accuracy vs computation time

³Computations times are calculated with computation on parallel work on 6 cores i7 cpu

Acknowledgments

Special thanks to :

- <http://machinelearningmastery.com> for sharing methods and tricks on data science, especially the method used in this report.
- <http://template latex.com> for sharing the LATEX template used here.

1. Code

A.1 Load Data

```
library(doMC)
registerDoMC(cores=6) # parallel computation

library(data.table)
library(xtable)

# Downloading files if they don't already exist
if(!file.exists("datas/trainDatas.csv")){
  download.file("https://d396gusza40orc.cloudfront.net/
    predmachlearn/pml-training.csv", "datas/trainDatas.csv")
}
if(!file.exists("datas/testDatas.csv")){
  download.file("https://d396gusza40orc.cloudfront.net/
    predmachlearn/pml-testing.csv", "datas/testDatas.csv")
}

datas <- fread("datas/trainDatas.csv", na.strings=c("NA","", "#DIV/0!"))
quizz <- fread("datas/testDatas.csv", na.strings=c("NA","", "#DIV/0!"))
```

A.2 Simplify dataset and deal with NAs

A.2.1 NA distribution

```
# remove unwanted columns
clean1 <- subset(datas, select = c(-1:-7))
# look at NA distribution in predictors
colNA <- apply(clean1, 2, is.na)
meanNA <- apply(colNA, 2, mean)
index <- 1:ncol(colNA)
meanNA <- cbind(meanNA, index)
meanNA <- data.frame(meanNA)

# Plot the distribution
par(mar=c(4,4,1,2))
hist(meanNA$meanNA, main="", breaks=75, col="blue",
  xlab="mean of NA")
```

A.2.2 NA mean

```
# Calculate mean of NA in predictors
naSums <- colSums(is.na(clean1))
dirty <- subset(clean1, select = (naSums!=0))
dirtyColNA <- apply(dirty, 2, is.na)
dirtyMeanNA <- apply(dirtyColNA, 2, mean)
resultsNA <- data.frame(mean=mean(dirtyMeanNA), sd=sd(dirtyMeanNA))
sumNA <- summary(dirtyMeanNA)
mini <- round(min(dirtyMeanNA), 4)
maxi <- round(max(dirtyMeanNA), 4)
moy <- round(mean(dirtyMeanNA), 4)
sdev <- round(sd(dirtyMeanNA), 4)
sumNA <- data.frame(min=mini, mean=moy, max=maxi, sd=sdev)
# Remove all predictors with NA from the dataset
# clean is not the reference dataset.
clean <- subset(clean1, select = (naSums==0))

# Print the summary in nice table
print.xtable(xtable(sumNA,
  caption="Summary for mean of NA in predictors
  with NA",
  digits=4, label="namean"), include.rownames=F)
```

A.3 Fitting a model

A.3.1 Functions to automate process

```
library(caret)
seed <- 27
# create quickview datasets
set.seed(seed)
quickTrain <- createDataPartition(clean$classe, p=0.05)[[1]]
quickview <- clean[quickTrain,]

# functions to automate model test
algo <- function(algo, datas, seed=27, metrics="Accuracy",
```

```
  preprocess=c("center","scale"), controlMeth="repeatedcv",
  controlNum=10, controlRep=3,...) {
  # function to fit one algorithm algo on datas, with possibility to choose
  # trainControl, metric, preProc, seed. Datas are centered and scaled by default.
  # add computation time for training at the end of the result.
  require(caret)
  set.seed(seed)
  # set trainControl
  control <- trainControl(method=controlMeth, number=controlNum,
    repeats=controlRep)
  # set the timer
  ptm <- proc.time()
  fitAlgo <- train(classe~, data=datas, method=algo, metric=metrics,
    preProc=preprocess, trControl=control,...)
  # record computation time
  timeFit <- proc.time()-ptm
  # store results in a list : x[[y]][[1]] give model,
  # x[[y]][[2]] give computation time
  list(fitAlgo, timeFit[[3]])
}

accuAlgo <- function(fitAlgo){
  # research of the best accuracy
  accuracy <- max(fitAlgo[[1]]$results$Accuracy)
  # research of SD for this accuracy
  indacc <- which.max(fitAlgo[[1]]$results$Accuracy)
  # accuracySD <- fitAlgo[[1]]$results$AccuracySD[indAcc]
  meth <- fitAlgo[[1]]$method
  # store results in data.frame
  res <- data.frame(Accuracy = accuracy*100, Time=fitAlgo[[2]], method=meth)
  row.names(res) <- fitAlgo[[1]]$modelInfo$label
  res
}

sampleAccu <- function(fitAlgo){
  # store all accuracies computed with the controlTrain in order to boxplot
  # and t.test them
  sampleaccu <- fitAlgo[[1]]$resample$Accuracy*100
  resu <- data.frame(sampleaccu)
  names(resu) <- fitAlgo[[1]]$modelInfo$label
  resu
}

multiAlgo <- function(multi, datas, seed=27, metrics="Accuracy",
  preprocess=c("center","scale"),
  controlMeth="repeatedcv", controlNum=10,
  controlRep=3){
  # function to test several model and store results in a list
  resultats <- list()
  for (i in multi){
    resultats[[match(i,multi)]] <- list(algo=datas, algo=i,
      seed=seed, metrics=metrics,
      preprocess=preprocess,
      controlMeth=controlMeth,
      controlNum=controlNum,
      controlRep=controlRep)
  }
  resultats
}

multiAccu <- function(multi){
  # function to present accuracy, sd and computation time in a ordered datafram
  # use accuAlgo function to retrieve info for each element of the list
  lapp <- lapply(multi, function(y) accuAlgo(y))
  # loop to store info in table
  leng <- length(multi)
  comput <- NULL
  for (i in 1:leng){
    comput <- rbind(comput, lapp[[i]])
  }
  # order the table on decreasing accuracy
  comput[order(comput$Accuracy, decreasing=T),]
```

A.3.2 Model quick review

```
modelAll <- c("bagFDA", "C5.0", "dnn", "earth", "gbm", "gpls", "hddda",
  "lda", "knn", "mda", "nb", "nnet", "parRF", "pam",
  "pda", "rff", "rpart", "snn", "svmLinear", "svmRadial",
  "treebag")

fitAll <- multiAlgo(modelAll, datas=quickview)
orderAllMethod <- as.character(multiAccu(fitAll)$method)

# print the reviewed model in table.
long <- length(fitAll)
model <- NULL
model <- sapply(fitAll, function(x) x[[1]]$modelInfo$label)
model <- data.frame(model)
```

```

print.xtable(xtable(model, caption="Models reviewed",
                     label="nameTable"), include.rownames=FALSE,
                     include.colnames=FALSE, hline.after=c(0, long),
                     table.placement="H")

# boxplot all model, ordered by accuracy
library(ggplot2)
plotsAll <- multiSampleAccu(fitAll)
plotsAccu <- multiAccu(fitAll)

library(reshape)
g <- ggplot(data=melt(plotsAll), aes(x=as.factor(variable), y=value))
g <- g + geom_boxplot() + theme_minimal()
g <- g + theme(axis.text.x = element_text(angle=30, hjust=1))
g <- g + ylab("Accuracy (%)"") + xlab("")
g

nameFive <- data.frame(row.names(plotsAccu[1:5]))
print.xtable(xtable(nameFive, include.rownames=F, include.colnames=F,
                     hline.after=NULL, table.placement="H"))

tTest <- apply(plotsAll[,1:5], 2, function(x) t.test(x, mu=mean(x)))
confInt <- sapply(tTest, function(x) x$conf.int)
confInt <- as.data.frame(confInt)
confTable <- matrix(nrow=5, ncol=2)
for (i in 1:5) {
  confTable[i,] <- confInt[,i]
}
row.names(confTable) <- names(plotsAll[,1:5])
print.xtable(xtable(confTable,
                     caption="Confidences intervals for models accuracy",
                     digits=4, label="confInt"), include.colnames=F,
                     hline.after=c(0,nrow(confTable)))

```

A.3.3 Best Five models

```

#bestFive <- c("C5.0", "rf", "parRF", "gbm", "treebag")
bestFive <- orderAllMethod[1:5]
fitFive <- multiAlgo(bestFive, datas=clean)

sampleAccuFive <- multiSampleAccu(fitFive)

g <- ggplot(data=melt(sampleAccuFive), aes(x=as.factor(variable),
                                             y=value))
g <- g + geom_boxplot() + theme_minimal()
g <- g + theme(axis.text.x = element_text(angle=30, hjust=1))
g <- g + ylab("Accuracy (%)"") + xlab("")
g

accuTimeFive <- subset(multiAccu(fitFive), select=c(-3))
accuTimeFive$Time <- round(accuTimeFive$Time, 0)
xtable(accuTimeFive, digits=2, label="tableFive",
       caption="Accuracy and computation time for the 5 bests models")

```

A.4 And the Winner is...

```

accuTimeFive$Accuracy <- accuTimeFive$Accuracy
attach(accuTimeFive)
par(mar=c(4, 4, 1, 1))
plot(Accuracy, Time, xlim=c(95,100.5), ylim=c(0,2300), pch=19, col=c(1:5),
      xlab="Model accuracy (%)", ylab="Computation time (s)")
text(Accuracy, row.names(accuTimeFive), pos=3)
detach(accuTimeFive)

```